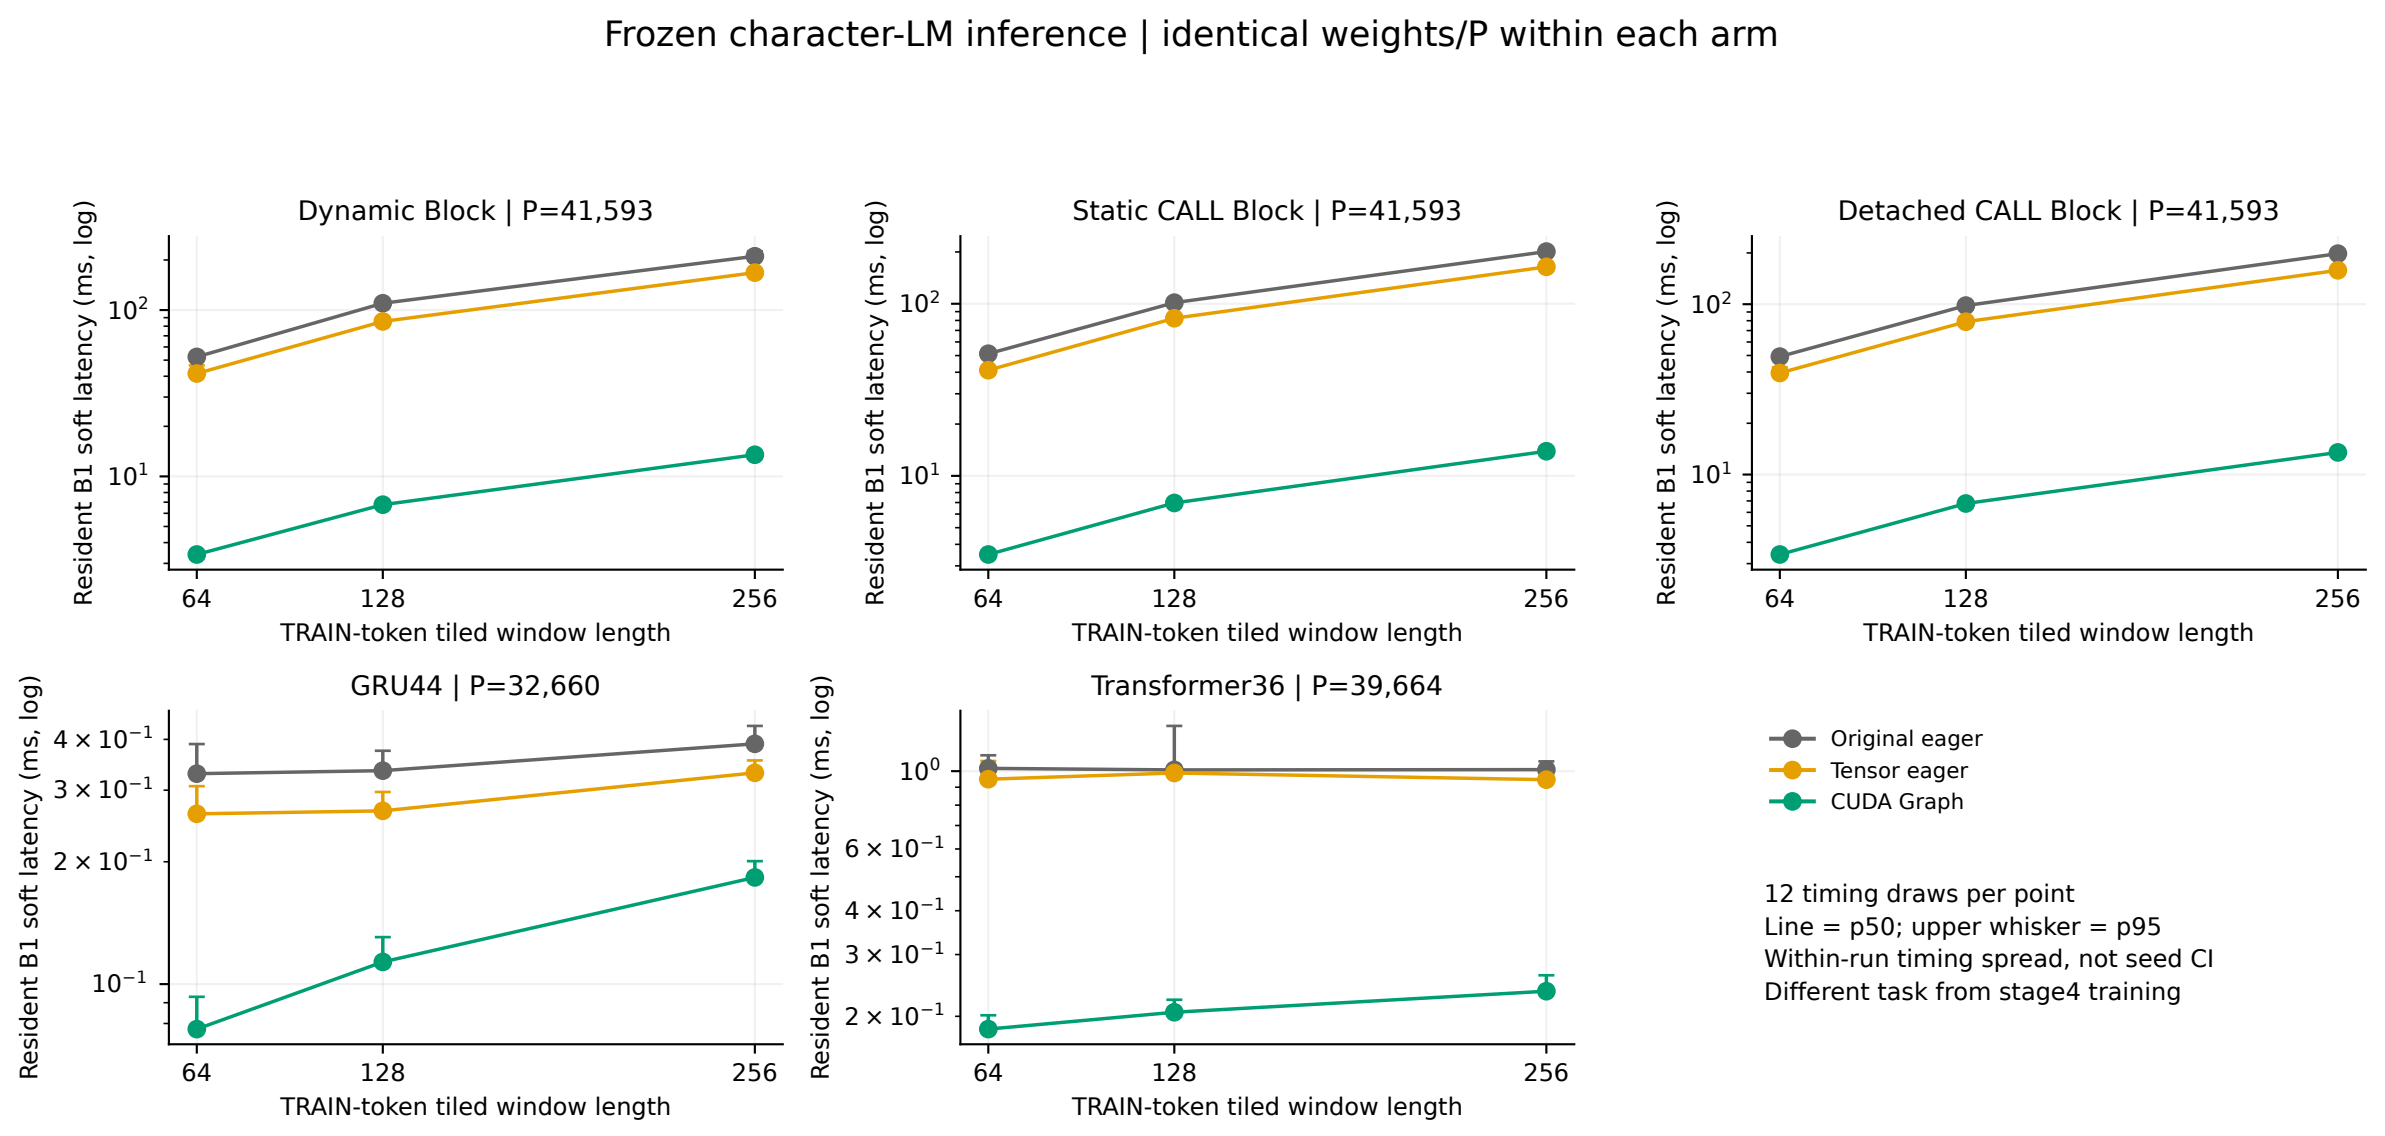

No training or labels in this frozen audit. B1, soft, first originalTRAIN windows tiled to64/128/256; outputs verified. Original-to-tensor removes checks AND all-active indexing; tensor-to-graph preserves the same tensor function. Resident-input only; excludes H2D/serving/cached sequential generation and makes no training-speed claim.